

Seminar

Haonan Lu

COMPASS

Feb 25, 2021

- ▶ Capture syscalls
 - > *Deploy sysdig*
 - > Write a kernel module manually to replace sysdig
- ▶ Make more experiments
 - > Figure out the overhead
 - > Judge whether the information captured is enough
 - > Make a demo to finish original schedule

- ▶ Capturing syscall: [Github](#)
 - > Finish a basic version (*print log, condition filter*)
 - > Investigate different types (i.e., passing different parameters) of syscalls
 - > Set this project as my *graduation project*

For example, we have a syscall `getrandom`:

```
ssize_t getrandom (void *__buffer, size_t __length, unsigned int __flags)
```

And there is the schema of syscall in Arm64:

Table 1: Syscall Calling Convention for Arm64

arch	syscall NR	return	arg0	arg1	arg2	arg3	arg4	arg5
arm64	x8	x0	x0	x1	x2	x3	x4	x5

Then we call it in our code:

```
int main(){  
    int a;  
    getrandom(&a, sizeof a, GRND_RANDOM);  
}
```

The capturing snippet show as follow: *but regs[0] was replaced by ret*

[my_sysdig:] syscall 0x116, with pid=0x178e, name=a.out

[my_sysdig:] exit syscall, regs[0]=0x4, with pid=0x178e, ret=0x4

There are 26 syscalls (arm64) similar to getrandom, and we need to save their regs[0] when entering the syscall.

- ▶ Set a global variable at entering and get its value at exit?
 - > No, multiple system calls can exist *at the same time*
 - > Considering multithread: When the program is waiting for a syscall (such as read), other threads can still run normally or call syscalls.
- ▶ Maintain a data struct to save regs[0]
 - > Syscall will block the thread, so we could save a value for each thread number.